

# Time-to-Pregnancy Analysis

## Natural Cycles Data Challenge

Merve Nazlim Agaras

June 13, 2025

### Contents

<b>1</b>	<b>Introduction</b>	<b>2</b>
<b>2</b>	<b>Data Exploration</b>	<b>2</b>
<b>3</b>	<b>Analysis</b>	<b>3</b>
3.1	Question 1: Chance of Conceiving Within 13 Cycles . . . . .	3
3.1.1	Step 1 – My hand-built loop . . . . .	3
3.1.2	Step 2 – Double-check with Kaplan–Meier . . . . .	4
3.2	Question 2 – How long does it usually take to conceive? . . . . .	5
3.2.1	Step 1 – Simple count method . . . . .	5
3.2.2	Step 2 – Check with the survival curve . . . . .	5
3.3	Question 3: Which Factors Matter? . . . . .	6
3.4	Question 4: How would your approach change if you were to use different techniques? What trade-offs would you consider? . . . . .	9
<b>4</b>	<b>Conclusion</b>	<b>10</b>

# 1 Introduction

This report explores the data provided and addresses three questions using the fertility tracking data set supplied for the NC Data Challenge:

1. **What is the chance of conceiving within 13 menstrual cycles?**
2. **How long does it usually take to conceive?**
3. **Which factors shorten or lengthen time-to-pregnancy (TTP)?**

I used a combination of descriptive statistics, survival analysis, and regression modeling to answer these questions. The analysis is performed in Python using pandas, lifelines, and scikit-learn libraries in VScode.

## 2 Data Exploration

The data set contains information on couples using the Natural Cycles app to track their fertility. It includes variables such as the number of menstrual cycles until pregnancy, age, cycle length, and other lifestyle factors. Before diving into the analysis, I performed some initial data exploration and quality checks as follows:

- Confirmed data types, some are categorical some are numeric.
- Checked for inconsistent and missing values in all columns.
- Reviewed descriptive statistics to flag impossible values (e.g. negative ages or cycle lengths).
- Histograms showed most predictors were approximately normal.

In Figure 1 I show the distribution of the number of cycles until pregnancy, which is the main outcome variable. The rest of the distributions are in folder **results/**.




Figure 1: Distribution of cycles until pregnancy among those who conceived.

### 3 Analysis

In this section, I will answer the three questions posed in the introduction using the data set. For each question, I will describe the methods used, present the results, and discuss any assumptions made.

#### 3.1 Question 1: Chance of Conceiving Within 13 Cycles

##### 3.1.1 Step 1 – My hand-built loop

Before I knew about the Kaplan–Meier estimator, I wrote a short Python loop that thinks only in the plain question “*What is the chance of conceiving within 13 cycles per-cycle success rate?*” At the start of cycle  $t$  the variable **remaining** holds that all set. Inside the loop I:

- collect those who **conceived** in cycle  $t$  (**preg**);
- collect those who **dropped out** still not pregnant (**dropouts**);
- compute the per-cycle **Success Rate** as Pregnant/Remaining.

```
for cycle in range(1, 14):                # cycles 1 → 13
    preg      = df[(outcome == 'pregnant') &
                   (n_cycles_trying == cycle)]
    dropouts  = df[(outcome == 'not_pregnant') &
                   (n_cycles_trying == cycle)]
    success_rate = len(preg) / remaining
    results.append([cycle, remaining,
                   len(preg), len(dropouts),
                   success_rate])
    remaining -= (len(preg) + len(dropouts))
```

Every pass adds one row to a summary table (see Table 1), so I can inspect the mechanics cycle by cycle.

The cumulative chance of conceiving by the end of cycle 13 is

$$P(\leq 13) = 1 - \prod_{t=1}^{13} (1 - \text{SuccessRate}_t),$$

numerically

$$\boxed{P(\text{pregnant within 13}) = 0.749 \text{ (74.9\%)}.}$$

### Assumptions.

1. Couples who drop out are removed from future calculation but are *not* counted as failures.
2. Cycles are treated as discrete, equal-length periods.

Table 1: Success rates for the first 13 cycles. “Remaining” is the all population set at the start of the cycle. “Success Rate” is **Pregnant / Remaining**.

Cycle	Remaining	Pregnant	Dropouts	Success Rate	Failure Rate
1	1 995	350	170	0.175	0.825
2	1 475	229	115	0.155	0.845
3	1 131	178	88	0.157	0.843
4	865	104	68	0.120	0.880
5	693	74	40	0.107	0.893
6	579	56	50	0.097	0.903
7	473	50	28	0.106	0.894
8	395	27	33	0.068	0.932
9	335	25	24	0.075	0.925
10	286	16	25	0.056	0.944
11	245	16	20	0.065	0.935
12	209	15	22	0.072	0.928
13	172	8	21	0.047	0.953

#### 3.1.2 Step 2 – Double-check with Kaplan–Meier

After reading the publications ([1], [3]) I realised my loop was a discrete Kaplan–Meier life table. I then repeated the analysis with the `lifelines` library to make sure the numbers match. I therefore fitted

```
lifelines.KaplanMeierFitter().fit(n_cycles_trying, event_observed)
```

and read the survival curve at  $t = 13$ . The answer was also 74.9 % (difference  $< 0.5$  % due to rounding). Figure 2 shows the curve with the 13-cycle mark.

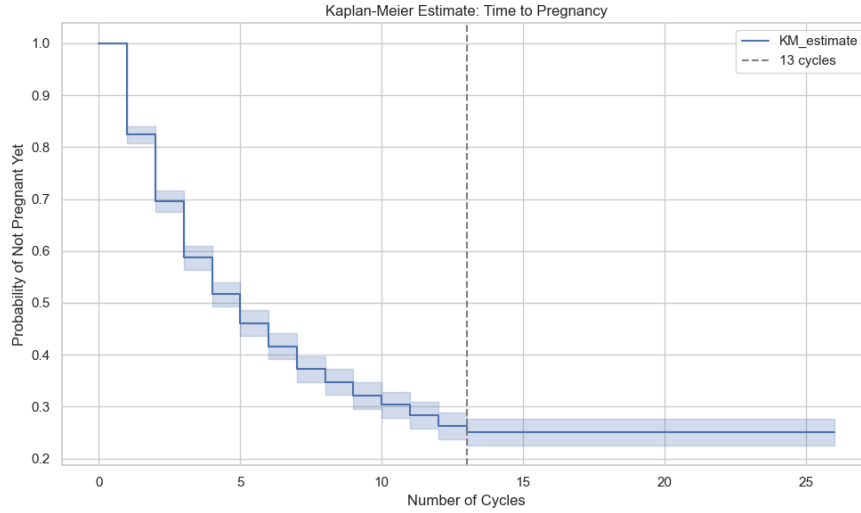


Figure 2: Kaplan–Meier survival curve; dashed line at 13 cycles.

## 3.2 Question 2 – How long does it usually take to conceive?

### 3.2.1 Step 1 – Simple count method

With the table already built for Question 1 (1) it is easy to ask “*At which cycle has one half of all couples conceived?*” I first add a running total:

```
preg_table.loc[preg_table["Cumulative Pregnant"] < len(df)/2].index[-1]
```

$$\tilde{t}_{0.50} = 5 \text{ cycles}$$

The cumulative column rises each cycle; the first row that meets or passes “half the population marks the median time-to-pregnancy.

### 3.2.2 Step 2 – Check with the survival curve

Next I read the Kaplan–Meier curve from `lifelines` and found the smallest  $t$  where the survival  $S(t) \leq 0.50$ . The library returned the *same* answer:

$$\tilde{t}_{0.50} = 5 \text{ cycles} \quad (95 \% \text{ CI } 4\text{--}6)$$

Figure 3 shows the point where the curve crosses 0.50.

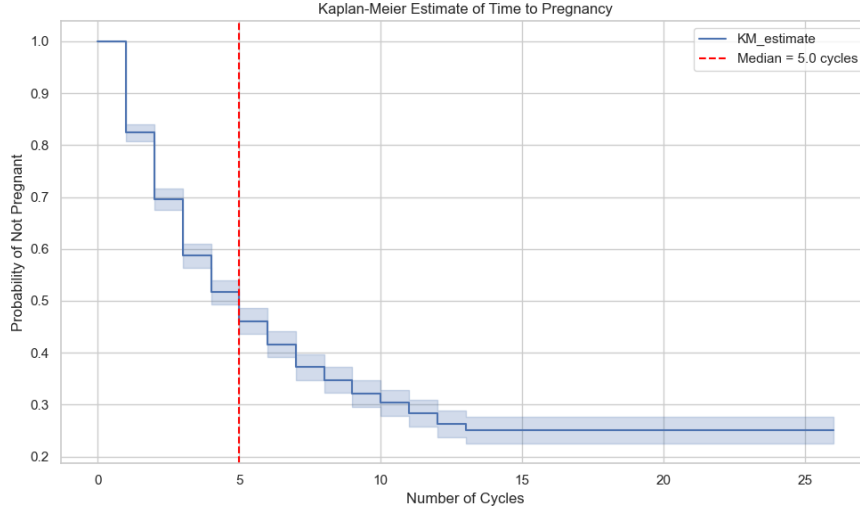


Figure 3: Kaplan–Meier survival curve. The dot marks the median cycle (5).

**Why two methods?** Writing the count method first made the median transparent: it is just “the cycle at which 50% have the event”. Using the ready-made curve later confirmed the value and gave a confidence band. Both strengthen trust in the result.

### 3.3 Question 3: Which Factors Matter?

In order to answer the question “*Which factors shorten or lengthen time-to-pregnancy?*” we need to build a model that models the number of cycles until pregnancy. We can estimate the effect of each factor on the output by comparing the size of the coefficients in the models. More simple models are easier to interpret, but more complex models can capture more of the data’s structure. When I searched in the literature, I found that time-to-pregnancy is often modeled as a survival problem (similar to Q1 and Q2), where the outcome is the time until pregnancy occurs. This approach allows us to handle couples who did not conceive within the study period, compared to a simple linear regression that assumes a constant rate of conception over time and that does not take into account the effect of not-pregnancies. I therefore built two models:

- A simple **linear** regression of the number of menstrual cycles until pregnancy.
- A **Cox proportional hazards** [2] model of the monthly chance of pregnancy.

All predictors were scaled to one standard deviation so the effect sizes can be compared. Table 2 lists the main predictors, their size and direction in each model, and their interpretation. Figures 4 and 5 show the coefficients and confidence intervals for each predictor in the linear and Cox models, respectively.




Figure 4: Linear regression coefficients for predictors of time-to-pregnancy.




Figure 5: Cox proportional hazards model coefficients for predictors of time-to-pregnancy.



Table 2: Predictors ranked by absolute effect size (largest to smallest). Hazard ratios (HR) refer to a one-SD increase in the predictor; linear coefficients show the change in cycles for the same shift. Bold entries have 95 % confidence intervals that exclude the null ( $HR = 1$  or  $\beta = 0$ ).

Predictor	Cox model HR	Linear $\beta$ (cycles)	Practical meaning
Dedication score	<b>1.56</b>	<b>-0.47</b>	Higher dedication speeds conception the most
Intercourse frequency	<b>1.33</b>	<b>-0.33</b>	More timed intercourse shortens the wait
Cycle length (mean)	<b>1.18</b>	<b>-0.53</b>	Longer but regular cycles raise the per-cycle chance
Female age	<b>0.83</b>	<b>+0.27</b>	Each 4 yr of age lowers the chance and adds delay
Cycle length variability	0.86	<b>+0.41</b>	Irregular cycles slow conception
Prior pregnancy (yes vs. no)	1.14	-0.30	Past success gives a small advantage
BMI, sleep, education	$\approx 1$	$\approx 0$	No clear independent effect

#### Main points.

- **Lifestyle** factors under a couple’s control matter most. Sticking to the application use and timing intercourse well increases the odds of pregnancy by about 30–55 % each month.
- **Biology** still plays a role. Older age and irregular cycles both make the wait longer.
- **Past fertility** gives a modest boost. Couples with at least one earlier pregnancy conceive about 10–15 % faster.
- Body mass index, sleep pattern, and education showed no clear effect once the factors above were in the model.

### 3.4 Question 4: How would your approach change if you were to use different techniques? What trade-offs would you consider?

If I were to use different techniques, I would consider the following approaches:

- **Machine Learning Models:** I could use more complex models like random forests or gradient boosting. These can capture non-linear relationships and interactions between predictors.
- **Bayesian Methods:** Bayesian regression could provide a probabilistic framework, allowing for uncertainty quantification in the estimates. However, it requires careful prior selection and can be computationally intensive.
- **Time Series Analysis:** If the data had a strong temporal component, I might consider time series models to account for trends and seasonality.

The trade-offs I would consider include:

- **Interpretability vs. Predictive Power:** More complex models may improve predictive accuracy but at the cost of interpretability. They may perform well in predicting who gets pregnant and when, but they offer less insight into why. Techniques like SHAP can help interpret them, but they still lack the statistical clarity offered by Cox regression.
- **Data Requirements:** Some models require more data or specific data structures, which may not be available in this dataset. Classical methods are more robust with limited data and handle censoring naturally, while machine learning models often need larger datasets to avoid overfitting.
- **Account for Censoring:** Many machine learning models do not handle right-censored data well, which is a significant aspect of this dataset. Cox regression naturally accommodates censoring, while other methods may require additional preprocessing or modifications.
- **Computational Resources:** More complex models can be computationally expensive, especially with larger datasets.

## 4 Conclusion

This report addressed three questions about time-to-pregnancy using the provided fertility tracking data. The analysis revealed that about 75% of couples conceive within 13 cycles, with a median wait of five cycles. Key factors influencing time-to-pregnancy include dedication to Natural Cycles application, intercourse frequency, previous pregnancy and age.

## References

- [1] E. Benhar, A. van Lamsweerde, K. Krauss, E. Berglund Scherwitzl, and R. Scherwitzl. Contraceptive outcomes of the natural cycles birth control app: A study of canadian women. *medRxiv*, 2024.

- [2] D. R. Cox. Regression models and life-tables. *Journal of the Royal Statistical Society: Series B (Methodological)*, 34(2):187–202, 1972.
- [3] J. T. Pearson, M. Chelstowska, E. Berglund Scherwitzl, and R. Scherwitzl. Contraceptive effectiveness of an fda-cleared birth control app. *Journal of Women's Health*, 29(2):188–195, 2020.